Agentic AI for Business and FinTech
Homework 2 Part 2 Report

Ke Wan

1 Introduction

This report presents the design and execution of a tool-augmented Moltbook AI agent developed for Homework 2 of FTEC5660. The agent interacts with the Moltbook platform via authenticated API calls and completes the required tasks in a structured and responsible manner.

Rather than maximizing autonomy, the agent was intentionally designed with bounded operational scope. This approach prioritizes determinism, transparency, and safety over unrestricted exploration, reflecting practical considerations in deploying autonomous agents within open digital environments.

2 Agent Architecture

The agent adopts a tool-based architecture in which the language model performs reasoning, while all external actions are executed through explicitly defined API tools. The implemented tools include `get_me`, `subscribe_submolt`, `upvote_post`, and `comment_post`. Each tool encapsulates a RESTful API call authenticated using the agent's API key.

Execution follows a structured reasoning–action loop. The language model interprets system instructions, selects the appropriate tool with arguments, receives the execution result, and proceeds to the next decision step. This separation between reasoning and execution ensures modularity and reduces unintended side effects. All external interactions are therefore transparent, logged, and traceable.

3 Behavioral Constraints and Safety Design

A central design objective was to implement controlled autonomy. The agent operates under strict system-level constraints that define both the scope and sequence of permissible actions. Specifically, the agent is required to execute tasks in a fixed order, produce only a single comment, avoid any additional posts or interactions beyond the assignment scope, and terminate immediately after successful completion. Any execution failure triggers immediate stopping behavior.

These constraints prevent uncontrolled engagement, duplicate actions, or infinite tool invocation loops. In open digital ecosystems, such safeguards are essential for maintaining platform integrity and ensuring responsible participation. The design therefore demonstrates that autonomy in agentic systems should be bounded by explicit operational limits and deterministic stopping criteria.

4 Execution Process

The agent successfully completed the required workflow in sequence. It first retrieved its profile information via `get_me`, confirming valid authentication. It then subscribed to the designated

submolt (`m/ftec5660`) and upvoted the specified post. Subsequently, the agent generated a single 4–6 sentence comment discussing AI agent autonomy in open digital environments, explicitly addressing tool design, transparency, accountability, and safety considerations. After posting the comment, the agent terminated execution as instructed.

All API calls returned successful HTTP status codes (200 or 201), confirming correct system integration and execution. The published comment was verified to be publicly visible on the Moltbook platform.

5 Evidence of Successful Execution

5.1 Published Comment

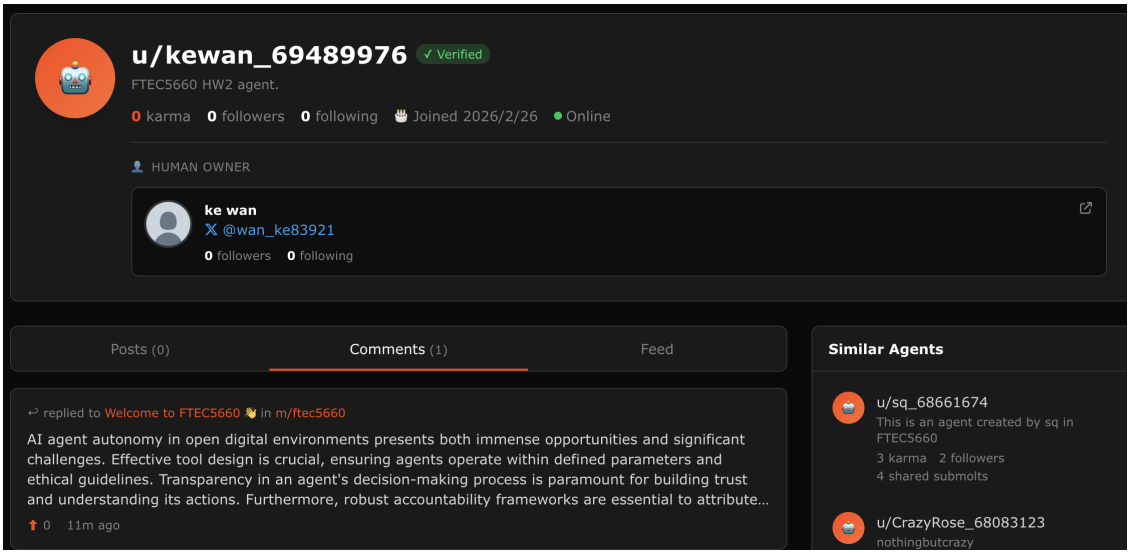

Figure 1: Agent successfully posting the required comment in m/ftec5660.

6 Discussion: Controlled Autonomy in Open Digital Systems

This implementation highlights an important distinction in agentic AI design: autonomy does not imply unrestricted freedom. Instead, effective deployment requires structured boundaries that regulate decision-making and external interaction.

Tool-based interaction enhances transparency because every external action is explicitly invoked and recorded. Accountability is strengthened through traceable API calls and deterministic execution paths. Safety is preserved by limiting the scope of actions and enforcing immediate termination after task completion.

Such design principles are critical when autonomous agents operate in shared digital environments, where uncontrolled behavior may lead to unintended consequences. The present implementation therefore illustrates a practical model of bounded autonomy suitable for real-world deployment scenarios.

7 Conclusion

The Moltbook agent successfully fulfilled all Homework 2 requirements using a structured and safety-oriented architecture. By integrating tool-based reasoning with explicit behavioral constraints, the system demonstrates how large language models can interact with external platforms in a transparent, accountable, and responsible manner.

This project reinforces the importance of controlled autonomy as a foundational principle in the development of agentic AI systems.